

Machine Learning L+Pr

Béla J. Szekeres, PhD

## Lecture 5

ELTE Faculty of Informatics, Szombathely, Hungary



## 1 Overview

## 2 Boosting

- Difference between Bagging and Boosting
- Boosting
- Math

## 3 XGBoost

## 4 Support Vector Machines

- Linearly separable case
- Math
- Problem
- Allowing misclassifications
- Nonlinear case

# Topics of this day

- Boosting
- Support Vector Machines
- Practice: Scikit-learn

## Difference between Bagging and Boosting

- it can be used for classification and regression too
- helps to **reduce variance and bias**

### What was **bagging**?

creates multiple copies of the original data: constructs several decision trees on the copies and combining all the trees to make predictions

**We construct these trees independently!**

### What is **boosting**?

the decision trees are grown sequentially so each tree is grown using information from previously grown trees → boosting is a sequential learning algorithm

**These trees are not independent of each other!**

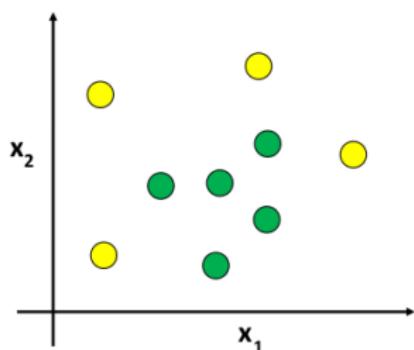
# Boosting

A rather counter-intuitive theory: a weak learner is not able to make good predictions

- weak learner is just a bit better than random guess or coin flip For example: decision trees with depth 1
- combining weak learners can prove to be an extremely powerful classifier
- by fitting small trees (decision stumps) we slowly improve the final result in cases when it does not perform well
- We will consider adaptive boosting "**AdaBoost**" algorithm

## Boosting: example

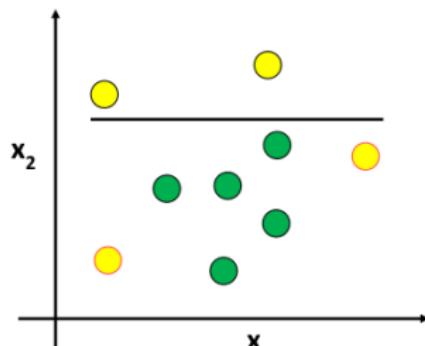
What is the task?



- we want to classify the dots
- two features
- + two output classes (yellow and green)
- **Boosting:** combines very simple weak learners such as decision trees with depth 1 (capable of linear classification)

## Boosting: example

What is the task?

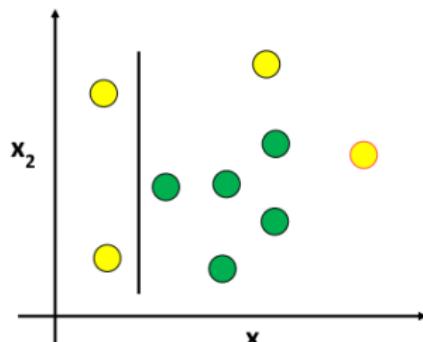


- we want to classify the dots
- two features
- + two output classes (yellow and green)
- **Boosting:** combines very simple weak learners such as decision trees with depth 1 (capable of linear classification)

- the classifier made 2 mistakes: two yellow dots are misclassified
- → in the next iteration it will focus on the misclassified items
- we'll increase the weights of the misclassified items and decrease the weights of the correctly classified items

## Boosting: example

What is the task?

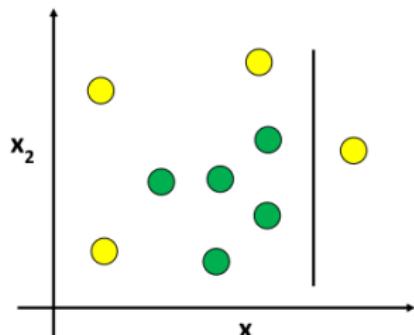


- we want to classify the dots
- two features
- + two output classes (yellow and green)
- **Boosting:** combines very simple weak learners such as decision trees with depth 1 (capable of linear classification)

- the classifier made 2 mistakes: two yellow dots are misclassified
- in the next iteration it will focus on the misclassified items
- we'll increase the weights of the misclassified items and decrease the weights of the correctly classified items

# Boosting: example

What is the task?

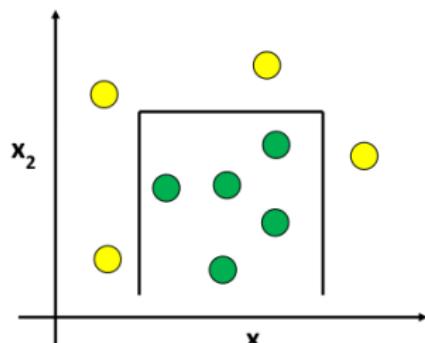


- we want to classify the dots
- two features
- + two output classes (yellow and green)
- **Boosting:** combines very simple weak learners such as decision trees with depth 1 (capable of linear classification)

- the classifier made 3 mistakes: three yellow dots are misclassified
- → in the next iteration it will focus on the misclassified items
- we'll increase the weights of the misclassified items and decrease the weights of the correctly classified items

## Boosting: example

We just have to combine these weak classifiers!



- we want to classify the dots
- two features
- + two output classes (yellow and green)
- **Boosting:** combines very simple weak learners such as decision trees with depth 1 (capable of linear classification)

- the classifier made 2 mistakes: two yellow dots are misclassified
- → in the next iteration it will focus on the misclassified items
- we'll increase the weights of the misclassified items and decrease the weights of the correctly classified items

## Math

- we keep combining  $h(x)$  weak learners (each learner knows just a little fraction of the space):  $H(x) = \text{sign} \sum_{i=1}^n \alpha_i h_i(x) \rightarrow$  Final model, not a weak learner anymore
- we assign  $+1$  and  $-1$  for the output classes (yellow and green)  $\rightarrow$  true class denoted by  $y$
- we initialize the (other) **weight parameters** at the beginning  $w_i = 1/N$
- $N$  is the number of the data
- make sure this is a **distribution**  $\sum_{i=1}^N w_i = 1$
- $\varepsilon = \sum_{\text{wrong}} w_i$ : the **error** is the sum of the misclassified weights

# Math: Algorithm

## The algorithm

initialize the weights  $w_i^1 = \frac{1}{N}$



pick  $h_t(x)$  tree that minimizes  
the  $\varepsilon_t$  error term

$$\varepsilon_t = \sum_{\text{wrong}} w_i^t$$



we can calculate  $\alpha_t$   $\alpha_t = \frac{1}{2} \log \left( \frac{1-\varepsilon_t}{\varepsilon_t} \right)$



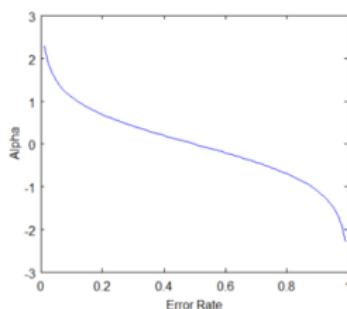
update  $w^{t+1}$  weights  $w_i^{t+1} = \frac{w_i^t}{Z} e^{-\alpha_t h_t(x_i) y(x_i)}$

Z:  $w^{t+1}$  should be a distribution

On every iteration we add a new  $h(x)$  weak learner to the final model

## Math: Algorithm

What is  $\alpha_t = \frac{1}{2} \log \left( \frac{1 - \varepsilon_t}{\varepsilon_t} \right)$ ? With given  $h(x)$  classifier...



- $\alpha$  increases as the error converges to 0:  
of course, good classifiers are given more weight
- $\alpha$  value is 0 if the error is 0.5.  
Why? Because it is a random guess (coin toss)  
→ we  
do not want our algorithm to rely on random guesses

- we give negative  $\alpha$  value for  $h(x)$  classifiers that are worse than random guesses  
⇒ we do the opposite that's the best action
- This  $\alpha$  parameter has something to do with the  $h(x)$  learners

## Math: Algorithm

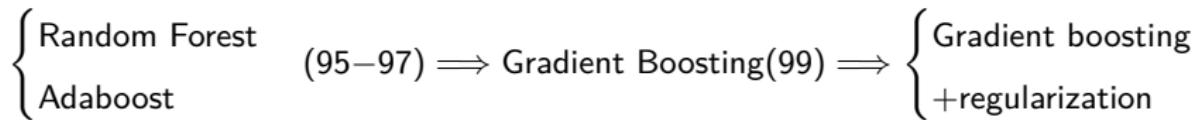
What? What is  $w_i^{t+1} = \frac{w_i^t}{Z} e^{-\alpha_t h_t(x_i) y(x_i)}$ ? ← weight updating

- the  $w$  weights have something to do with the dataset
- we set higher weights to more important samples and lower weight values to less important ones
- the  $Z$  makes sure  $w$  is distribution so the sum is 1
- $y(x)$  flips the sign of the exponent if  $h(x)$  is wrong  
Why is it good? It makes sure to assign smaller weights to samples that are correctly classified and bigger weights for misclassification  
→ in the next iteration the next  $h(x)$  learner can focus on those samples with higher weights
- Why to use  $\alpha$  in the formula? This is how we make sure that stronger classifiers' decisions are more important. If a weak classifier misclassifies an input we do not take that as seriously as a strong classifier's mistake

## Summary

Bagging	Boosting
every item has the same probability to appear in a new dataset	the samples are weighted so some of them will occur more often
parallel training stage	builds learners in sequential way
final decision is the average of the N learners	final decision is the weighted average of the N learners
reduces variance and solves overfitting	reduces bias but increases over-fitting a bit

# XGBoost



- eXtreme Gradient Boosted trees
- boosting ⊂ ensemble methods  $\implies$  Each tree boost attributes that led to misclassifications of previous tree
- **Tabular data**

## Advantages

- Routinely wins Kaggle competitions
- Easy to use, Fast
- High Model Performance
- A good choice for an algorithm to start with

# XGBoost

## Advantages - Features

- Regularized boosting  $\implies$  prevents overfitting
- can handle missing values automatically
- parallel processing, GPU
- tree pruning
- cross-validation at each iteration  $\implies$  enables early stopping
- sparse data

<https://arxiv.org/abs/1603.02754>

# Intuition - regression

## House price prediction

- Model 0 is the average of the target variable
- Average Price =  $(240 + 198 + 360 + 400)/4 = 299.5$

Squared meters	Price (thousands)	Residuals ( $y - \hat{y}_M_0$ )
120	240	-59.5
110	198	-101.5
200	360	60.5
400	400	100.5

# Intuition - regression

## House price prediction

- Model 1 is a decision tree using fitted on the residuals of Model-0

Squared meters	Price (thousands)	Residuals ( $y - \hat{y}_{M0}$ )	<b>Model 1 (predictions)</b>
120	240	-59.5	-30
110	198	-101.5	-90
200	360	60.5	50
400	400	100.5	75

# Intuition - regression

## Including the learning rate

Squared meters	Price (thousands)	Residuals ( $y - \hat{y}$ )	Model 1 (predictions)	M1 * lr + average
120	240	-59.5	$-30 * 0.5 = -15$	$299.5 - 15 = 284.5$
110	198	-101.5	$-90 * 0.5 = -45$	$299.5 - 45 = 254.5$
200	360	60.5	$50 * 0.5 = 25$	$299.5 + 25 = 349.5$
400	400	100.5	$75 * 0.5 = 37.5$	$299.5 + 75 = 374.5$

## Intuition - regression

Repeat the process multiple times

Average

$$+$$
 learning rate \* 

Model 0

$$+$$
 learning rate \* 

Model 1

$$+$$
 learning rate \* 

Model 2

Example with 2 trees:

$$299.5 + (0.5 * -15) + (0.5 * -30) = 277$$

# Hyperparameters

## Critical Hyperparameters

- *learning\_rate*: controls how much weight each tree has on the final prediction
- *num\_boost\_round*: defines the number of trees (number of iterations)
- *max\_depth*: controls the depth of the tree
- $\gamma$ : gamma - minimum require loss for the model to justify a split
- $\lambda$ : lambda - regularization on the weights
- *subsample*: percent of training data (rows) used for training a tree at each iteration
- *colsample\_bytree*: subsample ratio of columns when constructing each tree

$$\text{Loss} = \sum_{i=1}^n I(y_i, \hat{y}_i^{t-1} + f_t(x_t)) + \gamma T + \frac{1}{2} \lambda \sum_{j=1}^T \omega_j^2$$

# Howto use?

## When to use?

- large number of observations ( $>1000$ )
- less features ( $<100$ )
- performs well when data has mixture numerical and categorical features or just numeric features

## When don't use?

- computer vision,
- nlp

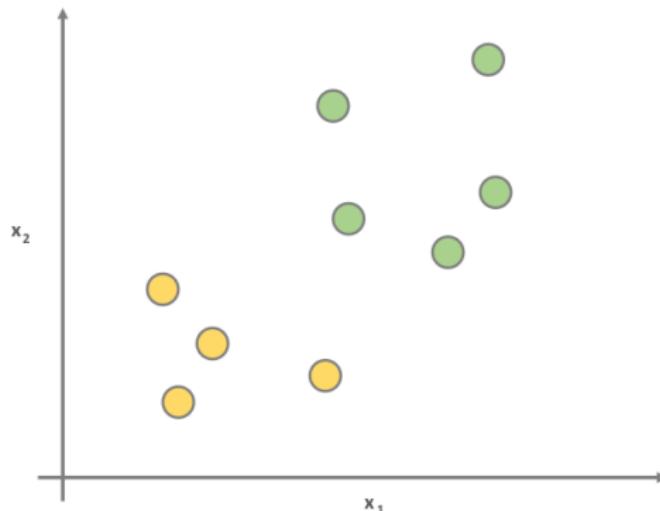
## Support Vector Machines

- very popular and widely used supervised learning algorithm (especially for **classification**)
- the great benefit is that it can **operate even in infinite dimensions**
- it defines a **margin** (decision boundary) between the data points in the multidimensional space
- that goal is to find a flat boundary (margin - optimal hyperplane in some sense) that leads to a homogeneous partition of the data
- a good separation is achieved by the hyperplane that has the largest distance to the nearest training-data point of any class since in general the larger the margin the lower the generalization error of the classifier

So with Support Vector Machine algorithm we maximize the margin.

# Support Vector Machines

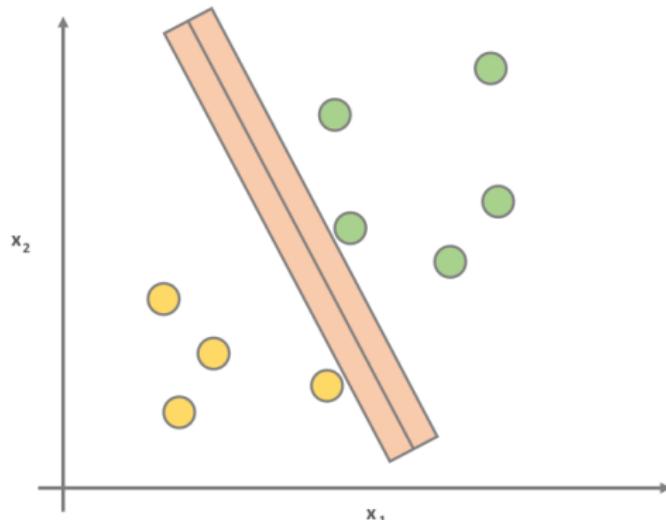
What is the aim?



- we want to find a **hyperplane**
- in this case a line that **separates** the data points with the **maximum margin**

# Support Vector Machines

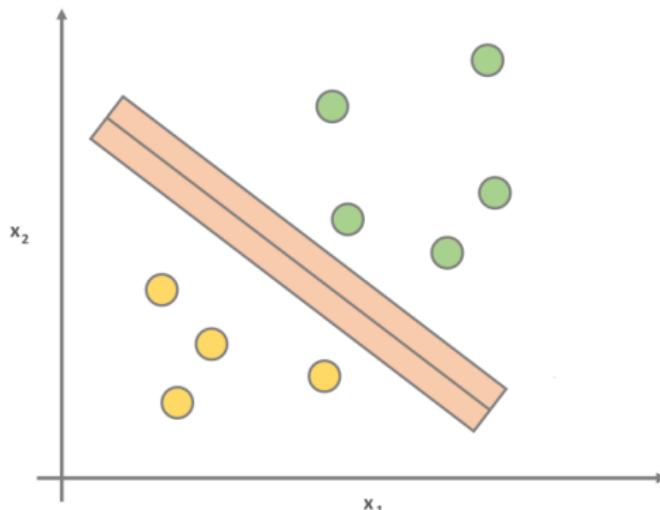
## What is the aim?



- we want to find a **hyperplane** - in this case a line that **separates** the data points with the **maximum margin**

# Support Vector Machines

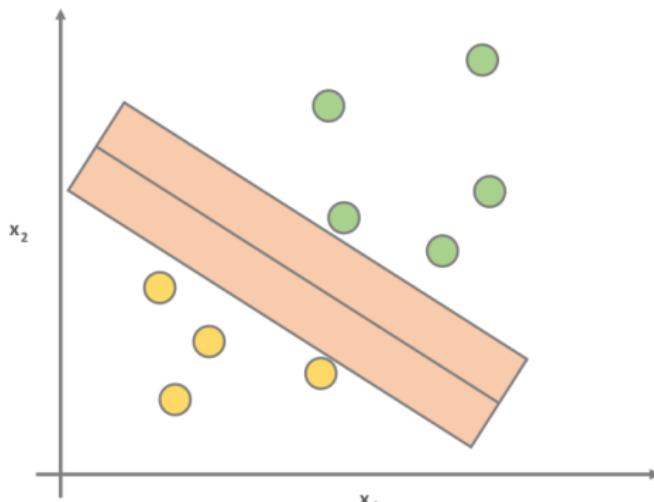
What is the aim?



- we want to find a **hyperplane**
- in this case a line that **separates** the data points with the **maximum margin**

# Support Vector Machines

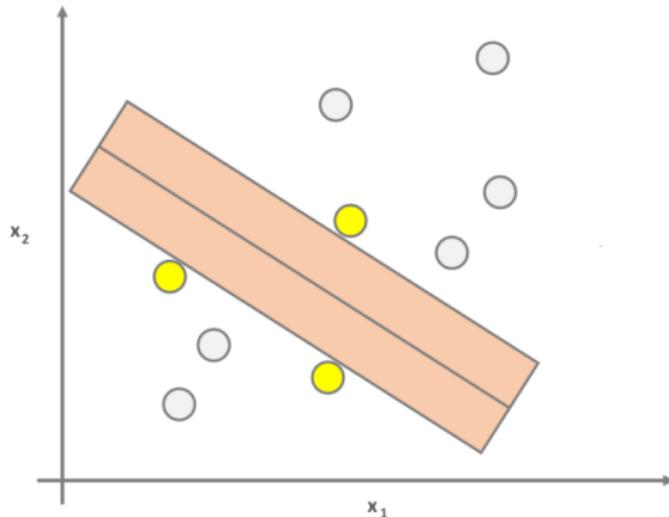
## The optimal margin



- we  
are looking for the margin  
(decision boundary) with  
the largest width  
**This is the maximum margin**

# Support Vector Machines

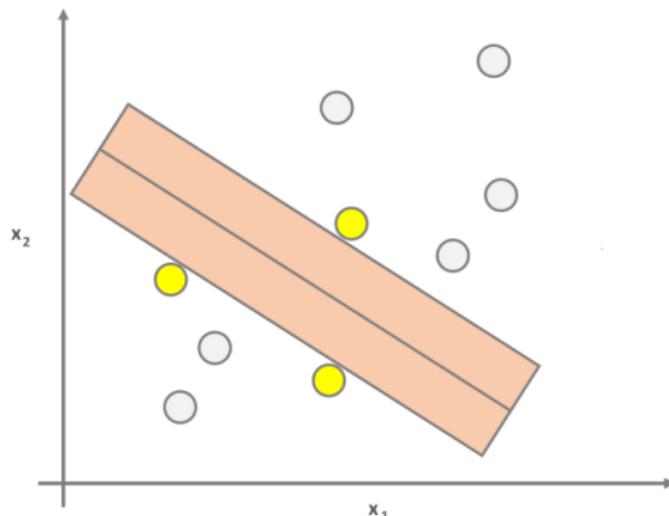
## The optimal margin: **support vectors**



- the points that are closest to the maximum margin are called the **support-vectors**
- storing the information for the decision boundary
- What if we change the other data points? Nothing

# Support Vector Machines

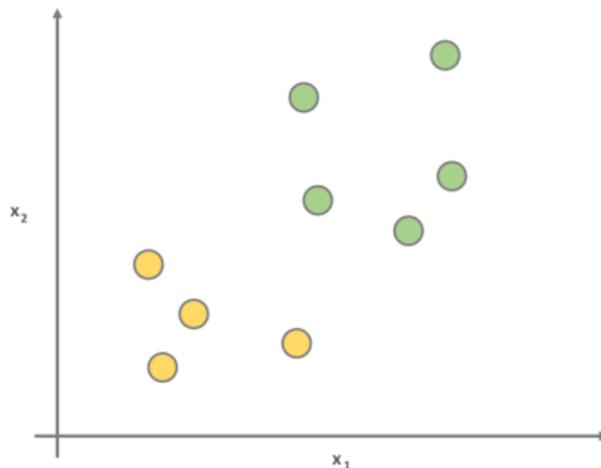
## The optimal margin: support vectors



- the points that are closest to the maximum margin are called the **support-vectors**
- storing the information for the decision boundary
- but if we change the support vectors  $\Rightarrow$  the model changes as well

## Math

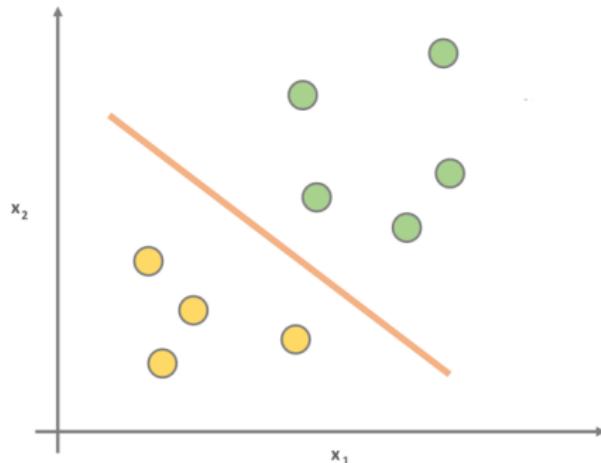
### The Hyperplane. What is a hyperplane?



in geometry a hyperplane is  
a subspace that has one dimension  
fewer than its ambient space  
- the hyperplane  
separates the space into two parts  
- the general equation  
for a hyperplane is  $\vec{w}^T \vec{x} + b = 0$

# Math

## The Hyperplane

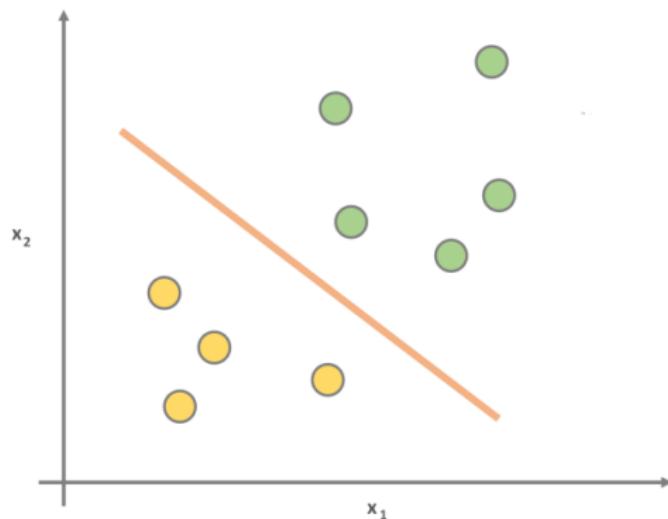


in geometry a hyperplane is  
a subspace that has one dimension  
fewer than its ambient space

- the **hyperplane**
- separates the space into two parts
- the general equation  
for a hyperplane is  $\vec{w}^T \vec{x} + b = 0$
- $x_2 = ax_1 + b$

## Math

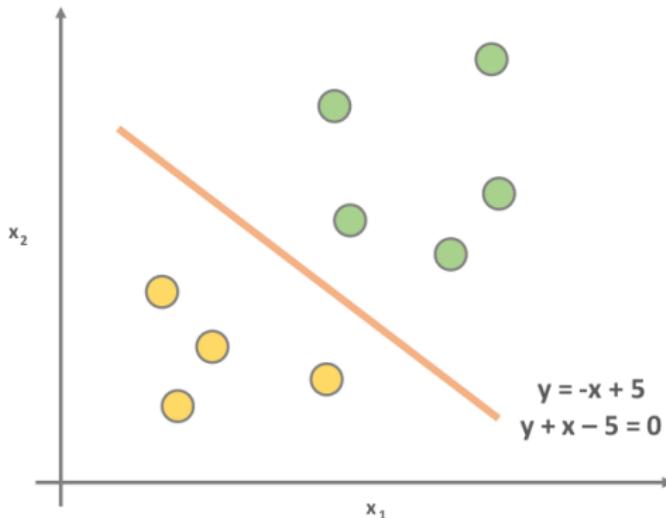
## The Hyperplane



in geometry  
a hyperplane is a subspace  
that has one dimension  
fewer than its ambient space  
- the hyperplane separates  
the space into two parts  
- the general equation for a  
hyperplane is  $\vec{w}^T \vec{x} + b = 0$   
-  $[1, -a] \cdot [y, x]^T - b = 0$

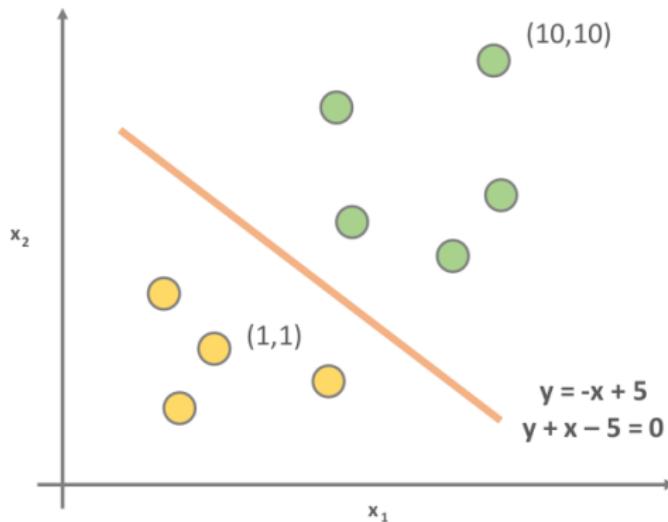
# Math: Example

## The Hyperplane



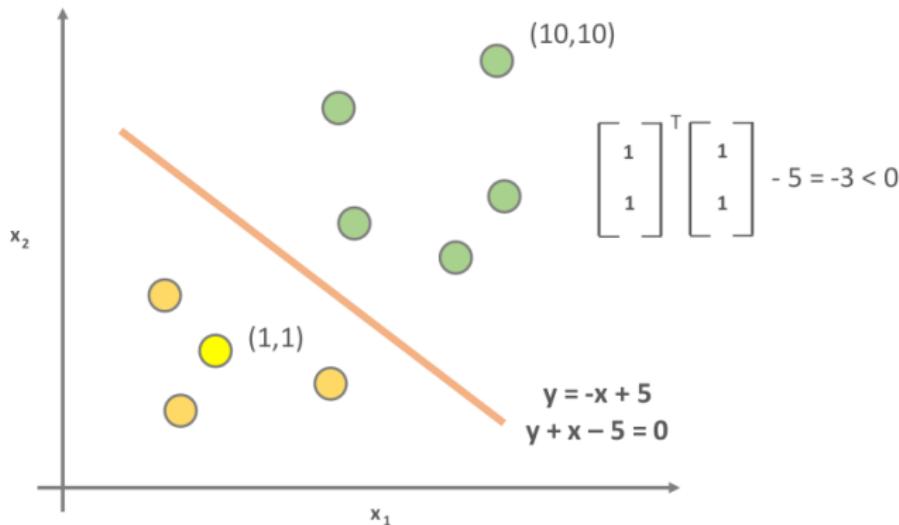
## Math: Example

## The Hyperplane



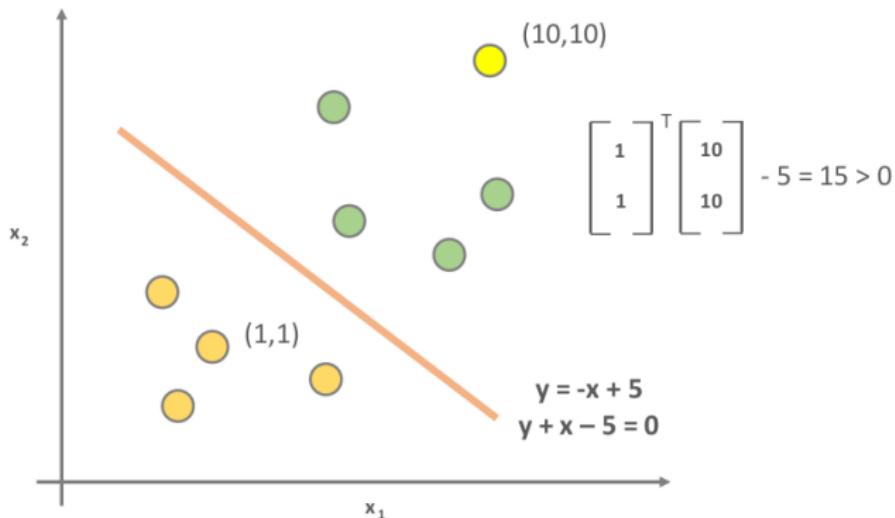
## Math: Example

## The Hyperplane



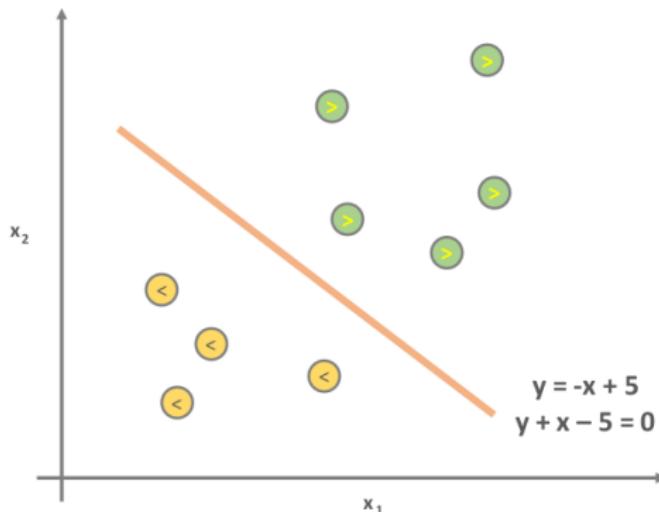
## Math: Example

## The Hyperplane



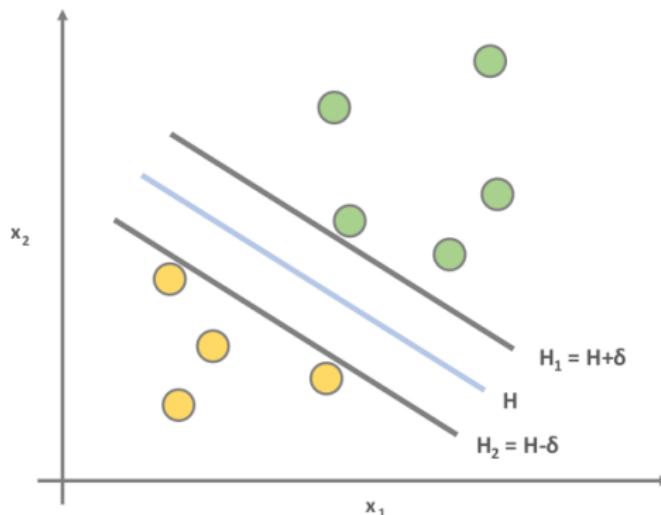
## Math: Example

## The Hyperplane



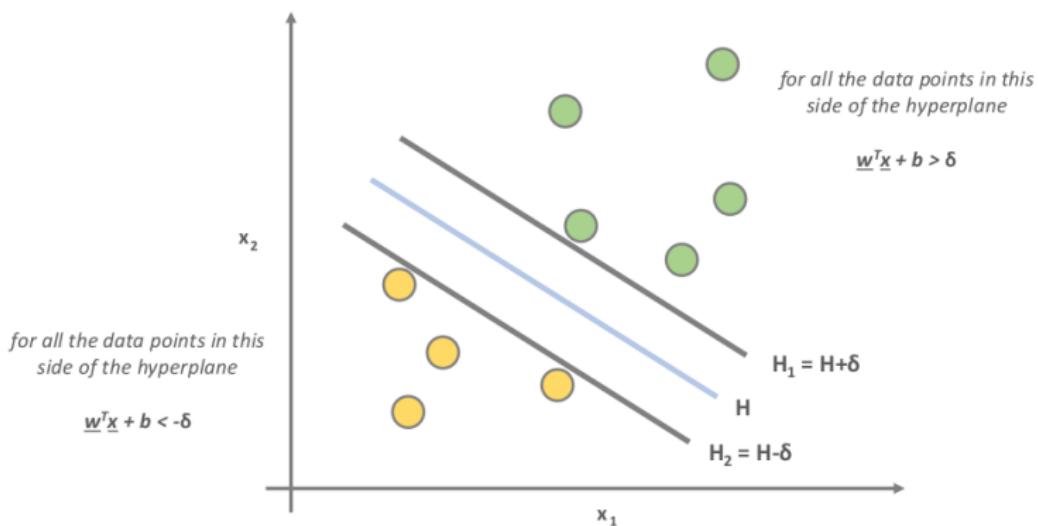
## Math: in general

## The Hyperplane



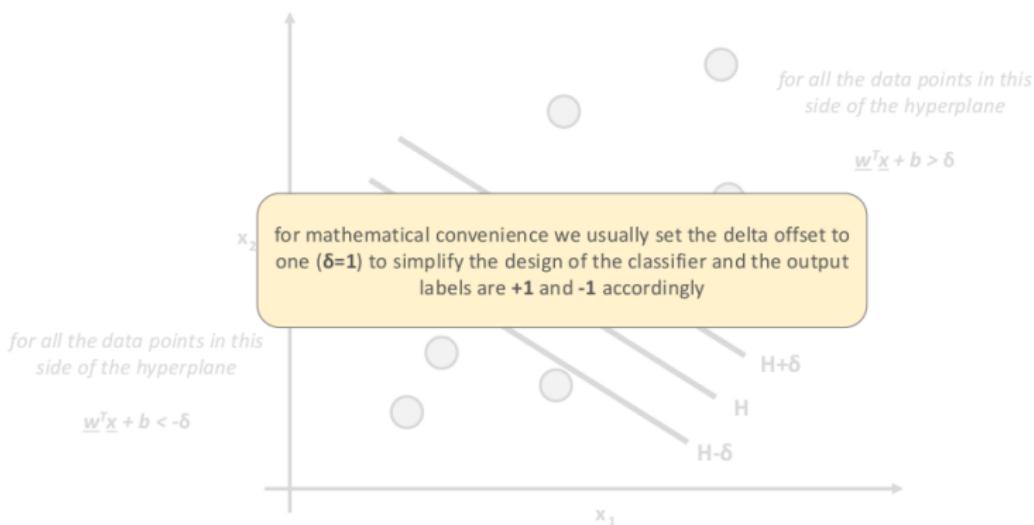
# Math: in general

## The Hyperplane



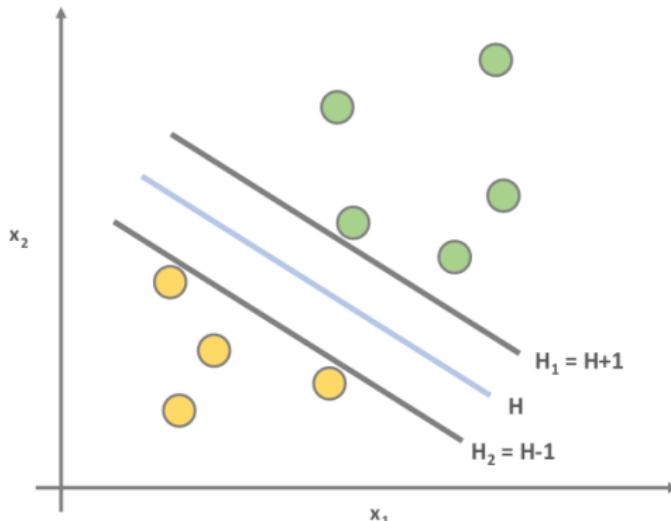
## Math: in general

## Important assumption



# Math: in general

## The constraints in one equation



*It is convenient to define  $\delta=1$  because it allows the classification constraints to be combined into a single constraint*

*we define the  $H$  hyperplane such that  
 $w \cdot x_i + b \geq 1$  when  $y_i = 1$  and  
 $w \cdot x_i + b \leq -1$  when  $y_i = -1$*

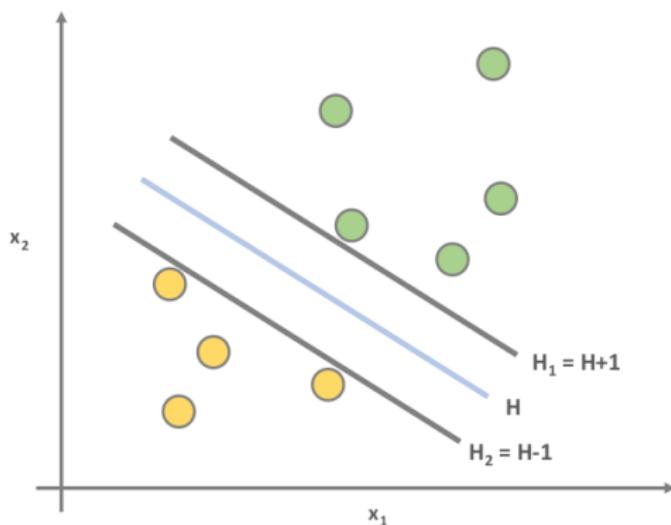
$$y_i(w \cdot x_i + b) \geq 1$$

*this is the single constraint we have to deal with !!!*

*(and of course support vectors are when  $y_i(w \cdot x_i + b) = 1$ )*

# Math: in general

## The distance



we define the  $H$  hyperplane such that  
 $\underline{w} \cdot x_i + b \geq 1$  when  $y_i = 1$  and  
 $\underline{w} \cdot x_i + b \leq -1$  when  $y_i = -1$

points on the  $H_1$  and  $H_2$  lines are called the support vectors

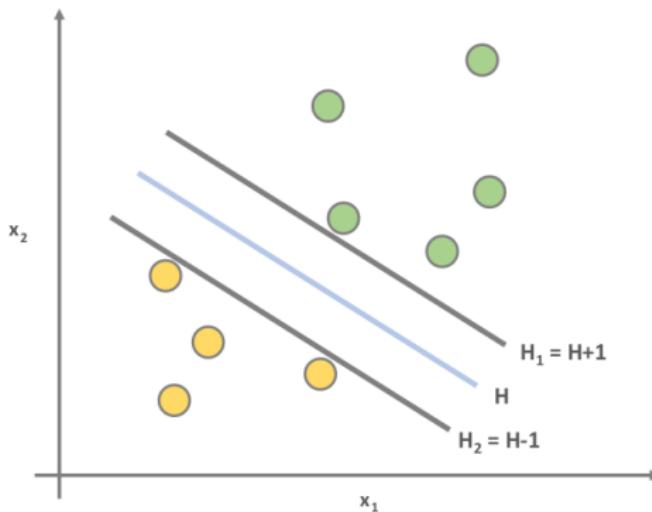
the distance between  $H_1$  and  $H_2$  lines is  $\delta_+ + \delta_-$

**WHAT IS THE TOTAL DISTANCE BETWEEN  $H_1$  AND  $H_2$ ?**

it is crucial because **SVM** is a maximum margin classifier so we have to **maximize the distance**

# Math: in general

## Maximization problem



the distance between  $H$  and  $H_1$  (and this is the distance between  $H$  and  $H_2$  as well) is defined by the following formula

$$\frac{1}{\|w\|}$$

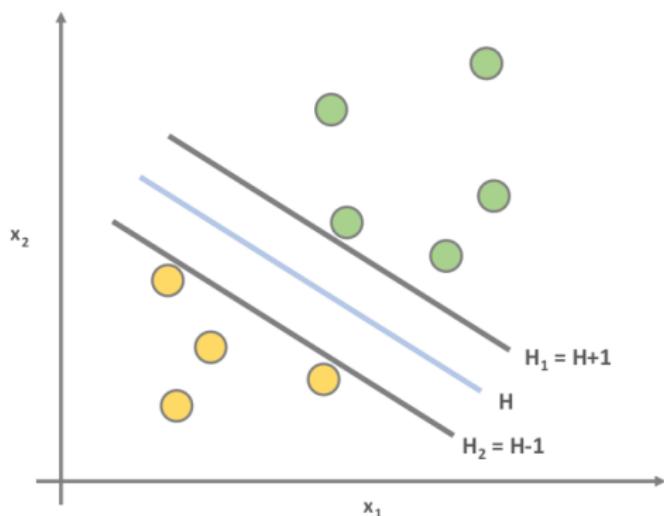
and if we want to get the total distance between  $H_1$  and  $H_1$  then:

$$\frac{2}{\|w\|}$$

so this is a typical **maximization problem**

# Math: in general

## Minimization problem



*every maximization problem can be transformed into a minimization problem*

$$\|w\|$$

*we can minimize this function instead and we can transform the problem into:*

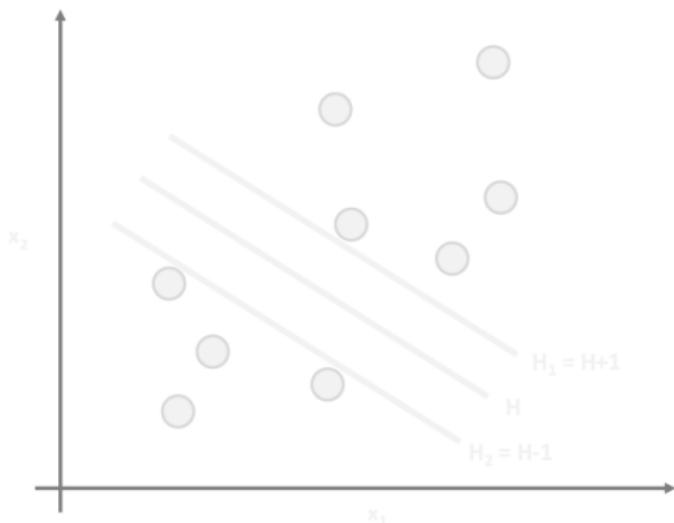
$$\min_w \frac{1}{2} \|w\|^2$$

$$y_i(\underline{w} x_i + b) \geq 1 \quad i=1\dots N$$

*so this is a typical minimization problem*

# Math: in general

## Summary



*every maximization problem can be transformed into a minimization problem*

$$||w||$$

### PRIMAL FORM OF

*we can minimize THE PROBLEM !!! instead and we can transform the problem into:*

$$\min_w \frac{1}{2} ||w||^2$$

$$y_i(\underline{w} x_i + b) \geq 1 \quad i=1 \dots N$$

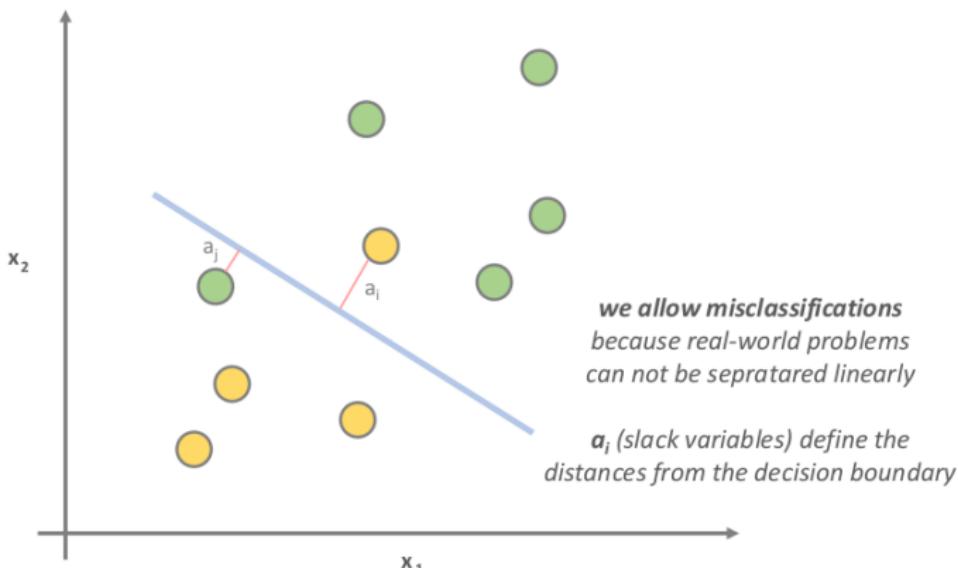
*so this is a typical minimization problem*

## Problem

- the main problem is that in real-world problems are non-linearly separable
- instead of classifying all the data points correctly - we **allow some mistakes**
- in many real-world applications the relationships between variables are usually non-linear
- a key feature of SVMs is their ability to map the problem into a higher dimensional space using a process known as the "kernel trick"
- ⇒ non-linear relationship may suddenly appear to be quite linear

# Allowing Misclassification

## Misclassifications



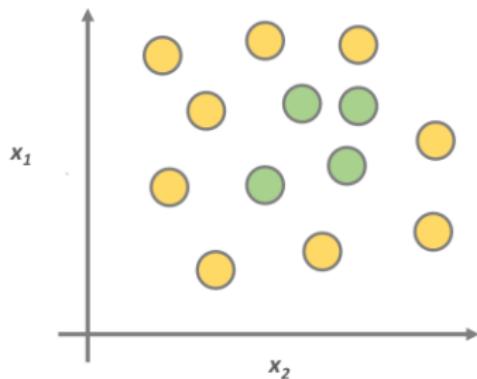
## Reformulation

### The problem statement

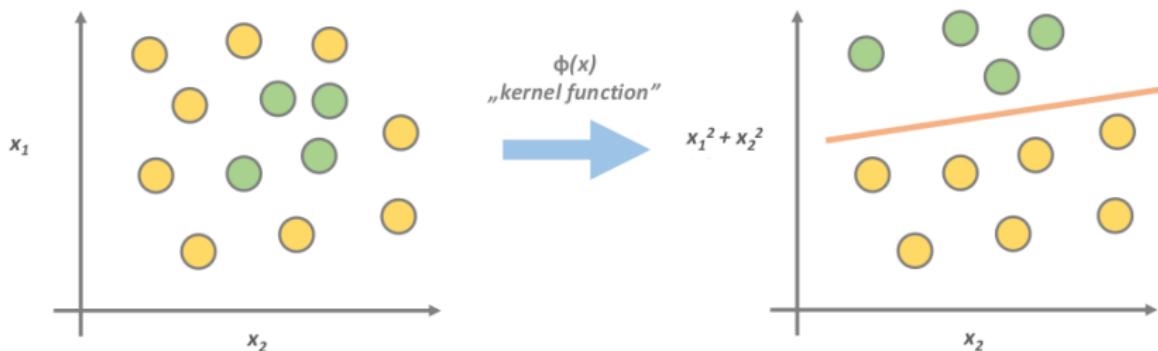
$$\begin{cases} \min_w \left[ \frac{1}{2} \|\vec{w}\|^2 + C \sum_i a_i \right] \\ y_i (\vec{w} \vec{x}_i + b) \geq 1 - a_i, \quad i = 1, \dots, N \end{cases}$$

- $C$  is the **Cost parameter** to all points that violate the constraints
- $a_i$  is 0 for the points that are classified correctly
- we can tune the  $C$  parameter - tuning the **Penalty term** for the data points that are misclassified
- if  $C$  is large then the algorithm tries to find a 100% separation
- if  $C$  is low then wider overall margin is allowed with more misclassified data points

## The kernel trick

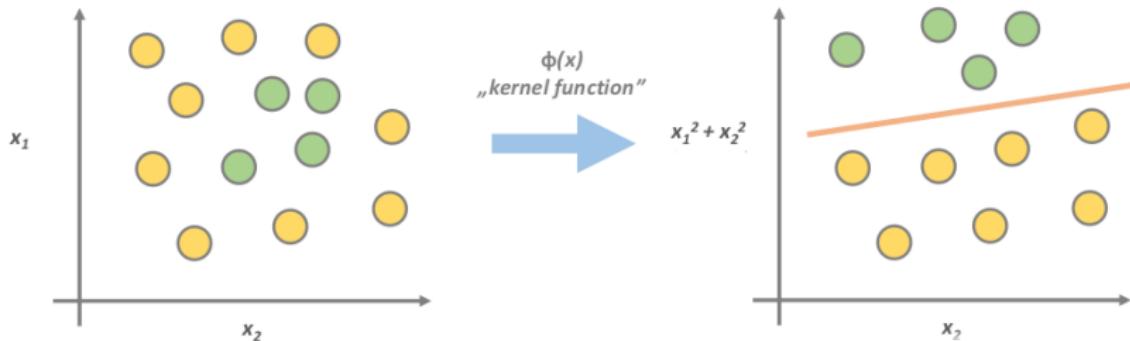


## The kernel trick



with the **kernel function** we can transform the problem into a higher dimensional space that is a linearly separable one (additional variable is altitude )

## The kernel trick



**SVM LEARNS CONCEPTS (FEATURE) THAT WERE NOT EXPLICITLY MEASURED IN THE ORIGINAL DATASET !!!**

## Summary

- so with the help of the  $\Phi(x)$  kernel-function we **transform all the points** in the dataset one by one
- and we end up with a **higher dimensional space**
- $K(\vec{x}_i, \vec{x}_j) = \Phi(\vec{x}_i)\Phi(\vec{x}_j)$

## Other Kernels

- linear kernel:  $K(\vec{x}_i, \vec{x}_j) = \vec{x}_i \cdot \vec{x}_j$ , it does not transform the data
- Polynomial** kernel:  $K(\vec{x}_i, \vec{x}_j) = (\vec{x}_i \cdot \vec{x}_j + 1)^d$
- Gaussian radial basis function** kernel:  $K(\vec{x}_i, \vec{x}_j) = \exp\left(-\frac{\|\vec{x}_i - \vec{x}_j\|^2}{2\sigma^2}\right)$

## Pros and contras

<b>Advantages</b>	<b>Disadvantages</b>
- can be used both regression and classification as well	- it deals with a large number of parameters and kernels
memory friendly	quite slow especially when there is a large number of features
it works fine even in infinite dimensions	there are no probabilities in the predictions