

# Machine Learning L+Pr

Béla J. Szekeres, PhD

Lecture 10

ELTE Faculty of Informatics, Szombathely, Hungary



## 1 Overview

## 2 Introduction

## 3 Gray-scaled images

- Kernel/Filter/Convolution
- Pooling
- Flattening

## 4 Convolutions on a color image

- Convolution over a volume
- Padding
- Stride
- Taking them together - Shape calculations

## 5 Example

# Topics of this day

- Motivation of Convolutional networks
- Basic building blocks
- Putting together

# Motivation

Problem definition: MNIST <http://yann.lecun.com/exdb/mnist/>

0	→	(1 0 0 0 0 0 0 0 0 0)
1	→	(0 1 0 0 0 0 0 0 0 0)
2	→	(0 0 1 0 0 0 0 0 0 0)
3	→	(0 0 0 1 0 0 0 0 0 0)
4	→	(0 0 0 0 1 0 0 0 0 0)
5	→	(0 0 0 0 0 1 0 0 0 0)
6	→	(0 0 0 0 0 0 1 0 0 0)
7	→	(0 0 0 0 0 0 0 1 0 0)
8	→	(0 0 0 0 0 0 0 0 1 0)
9	→	(0 0 0 0 0 0 0 0 0 1)


we're dealing  
with handwritten digit classification

- we have 10 classes
- represent the classes in binary
- **One-hot encoding**
- 60K training samples - 10K test
- $28 \times 28$  pictures

# Motivation

Problem definition: CIFAR-10

<https://www.cs.toronto.edu/~kriz/cifar.html>



we're dealing with  
a multiclass classification task again  
- we have 10 classes  
- 50K training samples - 10K test  
-  $32 \times 32 \times 3$  pictures  
→ we can  
transform it to a vector with size 3072

## Problem with deep networks

- Dense networks are working fine, but...
- if there are 1000 neurons in each layer, the number of weights increase dramatically
- the training will be extremely slow
- in the previous example for example with 1 hidden layer which contains 1000 neurons we have  $3072 \cdot 1000 + 1000 \cdot 10 + 10 + 1000 \approx 3M$  parameters
- and this was only a  $32 \times 32 \times 3$  shape picture

## Solution - Convolutional networks

- We know the inputs are images
- so we can encode certain properties into the network architecture

# Motivation

What is the main problem in machine learning? **Feature selection**

What features to use to build our model?



What features to use in order to recognize faces?

- location of nose
- distance between eyes



What features to use to recognize a cheetah?

- shape of ears
- black pattern

# Building blocks

We have 3 important steps

- convolutional operation (filter)
- pooling
- flattening

# Kernels

## Feature detector/kernel/filter


Feature detectors are represented as matrices

→ helps to detect the relevant features in a given image

0	-1	0
-1	5	-1
0	-1	0

This is the sharpen kernel

- increase the pixel intensity of the given one
- reduce it at the neighborhood



the original image



image after applying  
the sharpen kernel

# Kernels

## Feature detector/kernel/filter


Feature detectors are represented as matrices

→ helps to detect the relevant features in a given image

0	1	0
1	-4	1
0	1	0

This is the edge detection kernel

- decrease the pixel intensity of the given one
- don't change at the neighborhood



the original image




image after applying  
the edge detector kernel

# Kernels

## Feature detector/kernel/filter


Feature detectors are represented as matrices

→ helps to detect the relevant features in a given image

1	1	1
1	1	1
1	1	1

This is the blur kernel

- don't change the pixel intensity of the given one
- we use the neighborhood



the original image




image after applying  
the blur kernel

# Kernels

Every single kernel will use a specific feature of the image

When using edge detector, we assume the edges are important

## How to decide what feature detector to use?

- No need to decide in advance!
- convolutional networks use many kernels and during the training procedure it eventually select the best possible

# How does a kernel work?

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	0	1	1	0	0	0
0	0	0	1	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0

image



1	0	0
1	0	1
0	1	1

feature detector

=


feature map

# How does a kernel work?

0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	1	0	0	0	0
0	0	1	1	0	0	0	0
0	0	0	1	0	0	0	0
0	0	0	1	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0

image



1	0	0
1	0	1
0	1	1

feature detector

=

0			

feature map

## How does a kernel work?

0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	1	0	0	0	0
0	0	1	1	0	0	0	0
0	0	0	1	0	0	0	0
0	0	0	1	0	0	0	0
0	0	0	0	0	0	0	0



image

1	0	0
1	0	1
0	1	1

feature detector

=

0	1			

feature map

# How does a kernel work?

0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	1	0	0	0	0
0	0	1	1	0	0	0	0
0	0	0	1	0	0	0	0
0	0	0	1	0	0	0	0
0	0	0	0	0	0	0	0

image



1	0	0
1	0	1
0	1	1

feature detector

=

0	1	1	

feature map

# How does a kernel work?

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	0	1	1	0	0	0
0	0	0	1	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0

image



1	0	0
1	0	1
0	1	1


feature detector

=

0	1	1	0	

feature map

## How does a kernel work?



# How does a kernel work?

0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	1	0	0	0	0
0	0	1	1	0	0	0	0
0	0	0	1	0	0	0	0
0	0	0	1	0	0	0	0
0	0	0	0	0	0	0	0

image



1	0	0
1	0	1
0	1	1

feature detector

=

0	1	1	0	0
1				

feature map

# How does a kernel work?

0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	1	0	0	0	0
0	0	1	1	0	0	0	0
0	0	0	1	0	0	0	0
0	0	0	1	0	0	0	0
0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0

image



1	0	0
1	0	1
0	1	1

feature detector

=

0	1	1	0	0
1	3			

feature map

## How does a kernel work?

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	0	1	1	0	0	0
0	0	0	1	0	0	0
0	0	0	1	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0



image

1	0	0
1	0	1
0	1	1

feature detector

=


0	1	1	0	0
1	3	1	1	0
1	2	2	2	0
0	2	2	2	0
0	1	0	2	0

feature map

# How does a kernel work?

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	0	1	1	0	0	0
0	0	0	1	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0


image



# How does a kernel work?

0	0	0	0	0	0	0
0	0	0	0	0	0	0
0	0	0	1	0	0	0
0	0	1	1	0	0	0
0	0	0	1	0	0	0
0	0	0	1	0	0	0
0	0	0	0	0	0	0

image



## How does a kernel work?

### Spatial invariance

We would like to make sure to detect the same object no matter where it is located on the image or whether it is rotated/transformed




it is a cat: the location on the image does not matter or whether it is rotated or transformed

# Pooling

## Max pooling

With **max pooling** we select the most relevant features: this is how we deal with spatial invariance. We just care about the most relevant features.




- we can reduce the dimension of the image in order to end with a dataset containing the important pixel values without the unnecessary noise
- we reduce number of parameters: reduce overfitting

# Pooling

## Max pooling

With **max pooling** we select the most relevant features: this is how we deal with spatial invariance. We just care about the most relevant features.




- we can reduce the dimension of the image in order to end with a dataset containing the important pixel values without the unnecessary noise
- we reduce number of parameters: reduce overfitting

# Pooling

## Max pooling

With **max pooling** we select the most relevant features: this is how we deal with spatial invariance. We just care about the most relevant features.




- we can reduce the dimension of the image in order to end with a dataset containing the important pixel values without the unnecessary noise
- we reduce number of parameters: reduce overfitting

# Pooling

## Max pooling

With **max pooling** we select the most relevant features: this is how we deal with spatial invariance. We just care about the most relevant features.




- we can reduce the dimension of the image in order to end with a dataset containing the important pixel values without the unnecessary noise
- we reduce number of parameters: reduce overfitting

# Pooling

## Max pooling

With **max pooling** we select the most relevant features: this is how we deal with spatial invariance. We just care about the most relevant features.




- we can reduce the dimension of the image in order to end with a dataset containing the important pixel values without the unnecessary noise
- we reduce number of parameters: reduce overfitting

# Pooling

## Max pooling


With **max pooling** we select the most relevant features: this is how we deal with spatial invariance. We just care about the most relevant features.



- There are other techniques: **average pooling** is popular as well
- instead of choosing the maximum value we calculate the average of the values present in the subset


## Pooling

With max pooling we select the most relevant features: this is how we deal with spatial invariance. We just care about the most relevant features




## Flattening

The last operation we have to make is the flattening procedure: we transform the matrix into a one-dimensional vector: this is the input of a standard densely connected neural network



## Putting all together




# Convolutions on a color image - over a volume

## Tutorials


- [https://www.youtube.com/watch?v=KTB\\_0FoAQcc](https://www.youtube.com/watch?v=KTB_0FoAQcc)
- <https://gaussian37.github.io/dl-concept-cnn/>
- <https://guandi1995.github.io/One-Layer-of-a-Convolutional-Networks/>

## Convolutions on RGB image



## Convolutions on a color image - over a volume

Convolution = Scalar product over the volume




## Padding

Padding with zeros by  $P$  layers (containing 0s).

# Stride

## Pooling

Skipping  $S$  pixels.



## Pooling


You have to do the pooling by layers separately in spite of convolution.

# Stride


## Convolution

Skipping  $S$  pixels.

7 x 7 Input Volume



3 x 3 Output Volume



# Shape calculations

## Notations

- $f^{[l]}$ : filter size
- $p^{[l]}$ : padding
- $s^{[l]}$ : stride size
- $n_c^{[l-1]} = c$ : number of input channels
- $n_c^{[l]} = n_f$ : number of output channels = number of filters

## Input size

$$a^{[l-1]} = n_H^{[l-1]} \times n_W^{[l-1]} \times n_c^{[l-1]} = n_H^{[l-1]} \times n_W^{[l-1]} \times c$$

## Filter size

$$f^{[l]} \times f^{[l]} \times n_c^{[l-1]} = f^{[l]} \times f^{[l]} \times c$$

# Shape calculations

## Input size

$$a^{[l-1]} = n_H^{[l-1]} \times n_W^{[l-1]} \times n_c^{[l-1]} = n_H^{[l-1]} \times n_W^{[l-1]} \times c$$

## Filter size

$$f^{[l]} \times f^{[l]} \times n_c^{[l-1]} = f^{[l]} \times f^{[l]} \times c$$

## Output size


- $a^{[l]} = n_H^{[l]} \times n_W^{[l]} \times n_c^{[l]} = n_H^{[l]} \times n_W^{[l]} \times n_f$

- $n_H^{[l]} = \left[ \frac{n_H^{[l-1]} + 2p^{[l]} - f^{[l]}}{s^{[l]}} + 1 \right] \quad \text{and} \quad n_W^{[l]} = \left[ \frac{n_W^{[l-1]} + 2p^{[l]} - f^{[l]}}{s^{[l]}} + 1 \right]$

Kernel sizes and paddings can be different in each dimensions:

<https://pytorch.org/docs/stable/generated/torch.nn.Conv2d.html>

# Example: LeNet-5



## Other example

